

# Spell Checker Benchmark Report

---

*By Ivan Malishevskyi & Oleksii Savelich*

[All the code can be found here](#)

---

## Objective:

The goal of the task was to implement a simple spell checker that decides whether each individual word in a text file is spelled correctly, using various data structures. The evaluation was twofold:

- 1. Implement and test the spell-checking functionality using different approaches.
- 2. Compare running times for **dictionary building** and **spell checking** on a large piece of text.

## Implemented Approaches:

Four spell-checking strategies were implemented, each relying on a different underlying data structure:

Method	Description
LinearCheck	Naive linear search over a list of dictionary words
BstCheck	Uses a balanced binary search tree (e.g., SortedSet )
HashSetCheck	Employs a hash-based set for constant-time lookups
TrieCheck	Implements a trie (prefix tree) to match words by traversing character paths

---

## Experimental Setup

- **Dictionary Source:** A standard English word list (~100,000 words).

- **Text Sources:** Three input texts different in size, even though task suggested using "large piece of text", it wasn't specified what exactly word "large" means, so I thought I might use a couple of relatively large texts just in case.

1. `small.txt` (~10 KB)\*
2. `medium.txt` (~100 KB)\*
3. `AliceInWonderland.txt` (~150 KB)

\*marked files were made through modifying `AliceInWonderland.txt` by striping different amount of symbols at the end of the file


Each spell checker was benchmarked for:

1. **Dictionary build time** (structure construction from word list)
2. **Spell-check time** (checking a text file for misspelled words)

---

## Visual benchmark results followed by observations

Dictionary build time




### Observation:

HashSet is the fastest structure to build, followed closely by Linear. The Trie takes the

longest to build, likely due to its per-character node structure. BST sits in the middle but is significantly slower than HashSet.

Spellcheck time



## Spell-Check Time vs Text Size

Method	10 KB (ms)	100 KB (ms)	150 KB (ms)	Time Complexity
LinearCheck	460.79	3563.11	5389.57	$O(N \times D)$ — linear scan through the dictionary
BstCheck	15.57	134.85	203.71	$O(N \times \log D)$ — binary search in balanced BST
HashSetCheck	0.72	6.39	9.73	$O(N)$ — average-case constant lookup
TrieCheck	1.64	12.49	18.83	$O(N \times L)$ — prefix tree traversal per word

### Legend:

<b>N</b> = number of words in the input text
<b>D</b> = number of dictionary words
<b>L</b> = average word length

## Summary of Results

- **LinearCheck** is the slowest and scales poorly — shouldn't be used for large inputs.
- **BstCheck** improves performance with logarithmic lookups but still grows with input size.
- **HashSetCheck** is the fastest overall — great build time and constant-time lookups.
- **TrieCheck** performs nearly as well and is useful for prefix-based operations.

## Conclusion

This project clearly demonstrates how the choice of data structure directly impacts both the speed of building a dictionary and the efficiency of spell checking. The differences were especially noticeable when processing larger texts. But on a smaller scale texts it could be clearly seen too.

In summary, **hash-based and trie-based approaches should be preferred** for building high-performance spell checkers. They scale well with large texts and provide fast, predictable behavior.

---

#### Tools used: >

- [Obsidian](#) , *to generate this fancy report and export it into pdf*
  - [BenchmarkDotNet](#) , *to measure closest to raw performance of the algorithms*
  - [C# & .NET](#) , *the language of choice and environment to implement algorithms*
  - [Project Gutenberg library](#) , *the input text was downloaded from there*
  - [Matplotlib](#) , *library used for visualization output of benchmarks*
- 

#### Machine specs >

Windows 11 (10.0.22631.5189/23H2/2023Update/SunValley3)  
12th Gen Intel Core i5-1240P, 1 CPU, 16 logical and 12 physical cores  
.NET SDK 8.0.408  
[Host] : .NET 8.0.15 (8.0.1525.16413), X64 RyuJIT AVX2  
DefaultJob : .NET 8.0.15 (8.0.1525.16413), X64 RyuJIT AVX2

## Easter egg

I'm sure that nobody's gonna read my report till the end but in case you will I just wanna let you know that I love pizza ❤️